

# Word2vec

- word2vec architecture
- subsampling
- negative sampling
- GLOVE
- FastText
- ELMO

Seminar - ranking of messages

**Lecture 1. Milana Shkhanukova, 2023**

# Recap

- В чем был главный недостаток TF-IDF и One-Hot encoding?


# N-gram embeddings, CBOW

**Problem?** Our words appear together, it matters.

Guess the word by its context, CBOW – **Continuous Bag-of-Words**


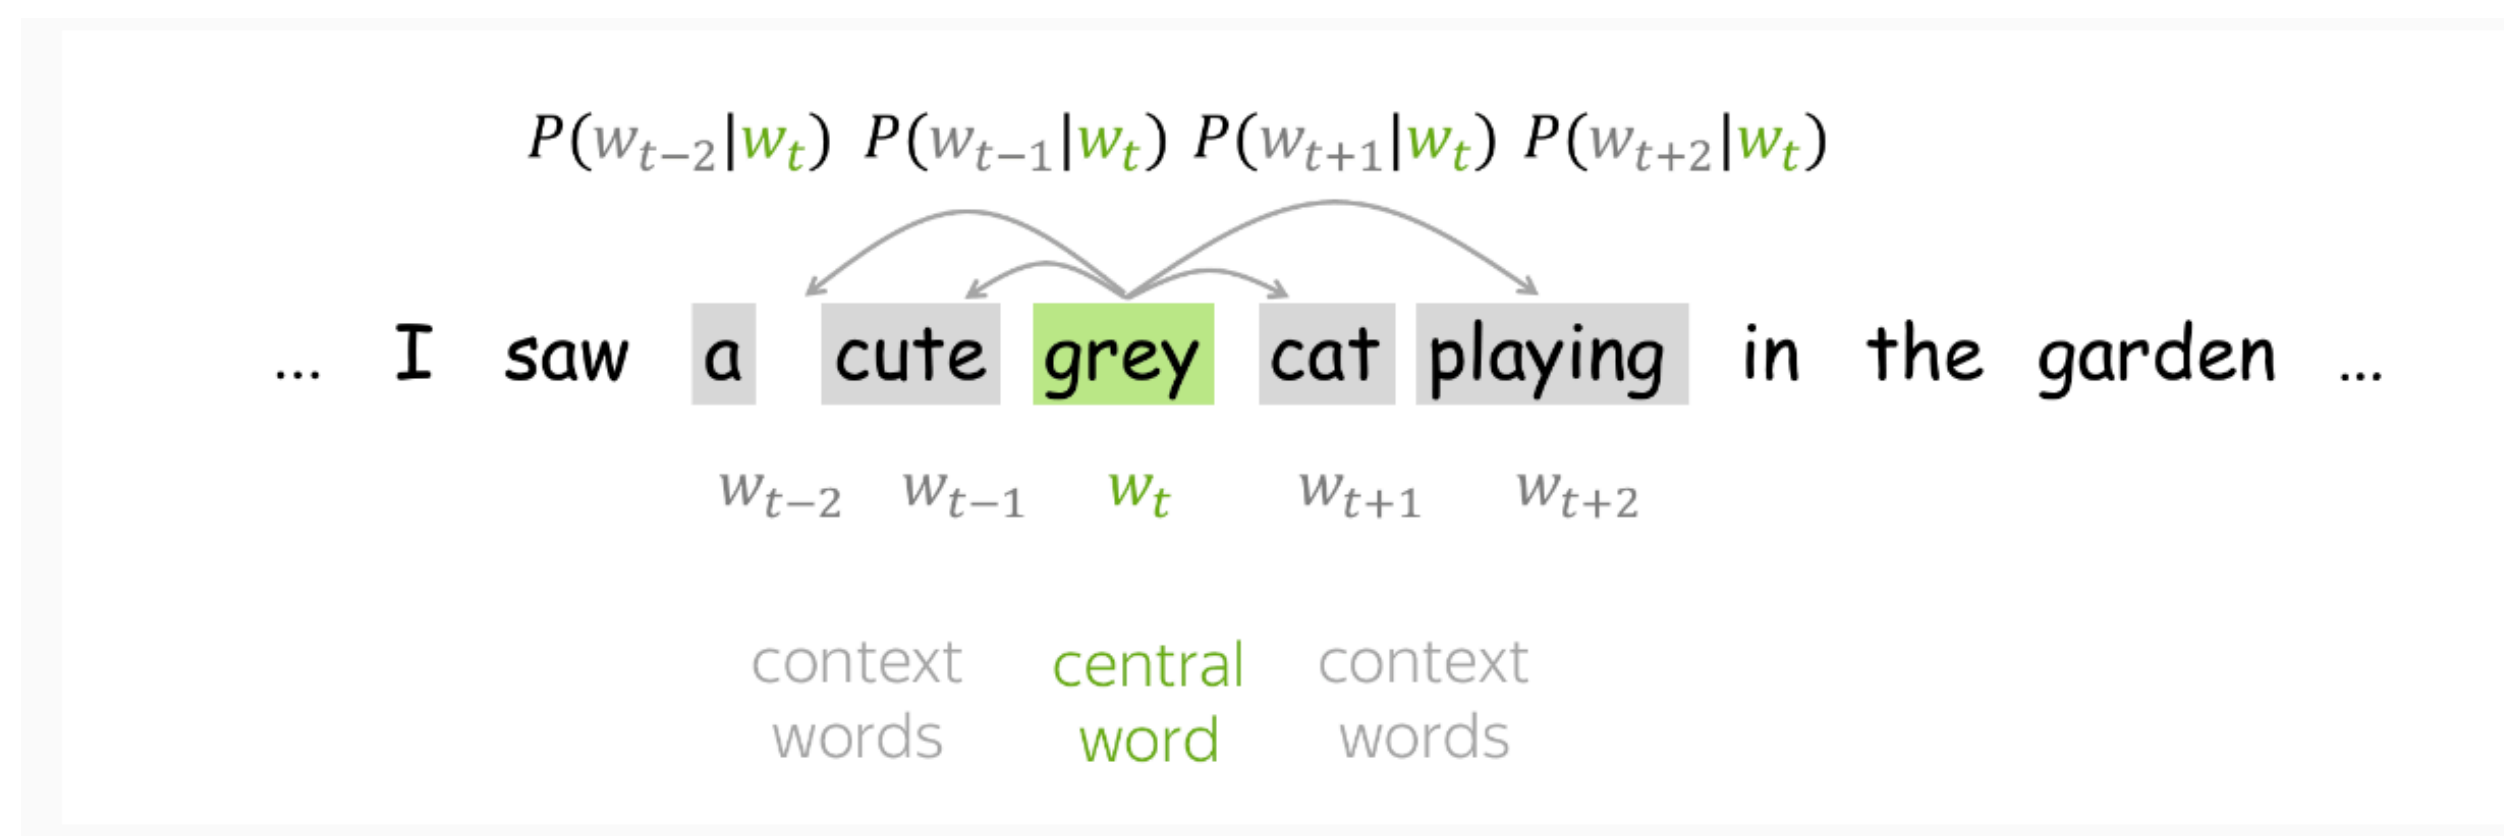
**Continuous** - continuous vector representations to represent words.  
Each word is represented as a dense vector in a continuous vector space.

Окно - это количество токенов в одну сторону, на которое мы смотрим



# N-gram embeddings, Skip-gram

Problem? Our words appear together, it matters.



Skip-gram

# What is n?

Я люблю пить кофе по утрам одна      Пить люблю я кофе по утрам одна

$N = 3$

Контекст = пить кофе по

Контекст = я люблю пить + по утрам одна

# N-window size

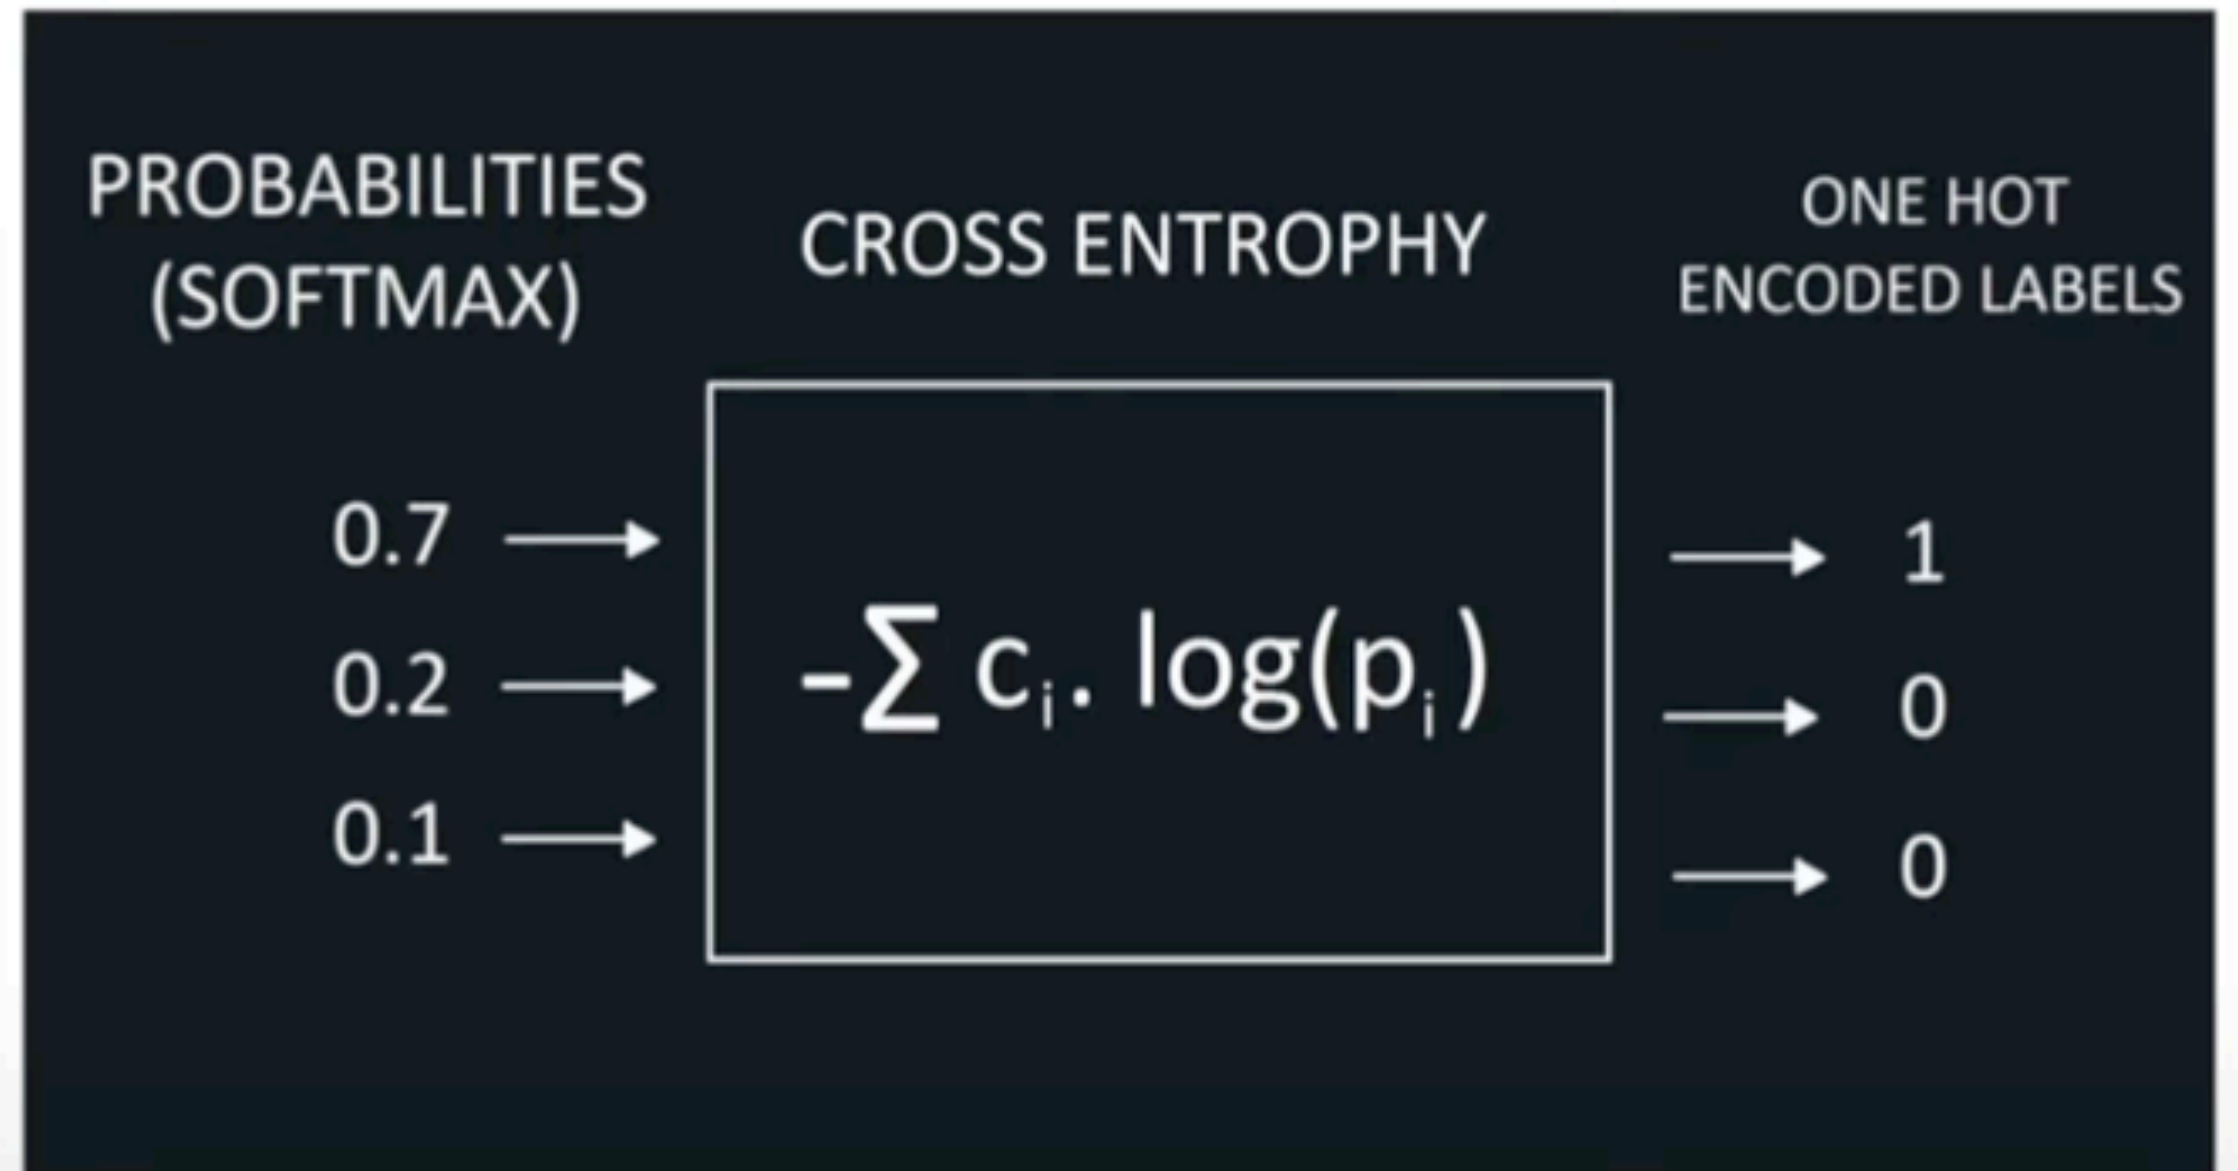
[this paper](#) notes that larger windows tend to produce more topical similarities (i.e. **dog**, **bark** and **leash** will be grouped together, as well as **walked**, **run** and **walking**), while smaller windows tend to produce more functional and syntactic similarities (i.e. **Poodle**, **Pitbull**, **Rottweiler**, or **walking**, **running**, **approaching**)

# Some questions

- Are context words at different distances equally important? If not, how can we modify co-occurrence counts?  
Контекст везде одинаковый?
- В какой модели у нас есть информация о порядке слов?

# Cross entropy loss

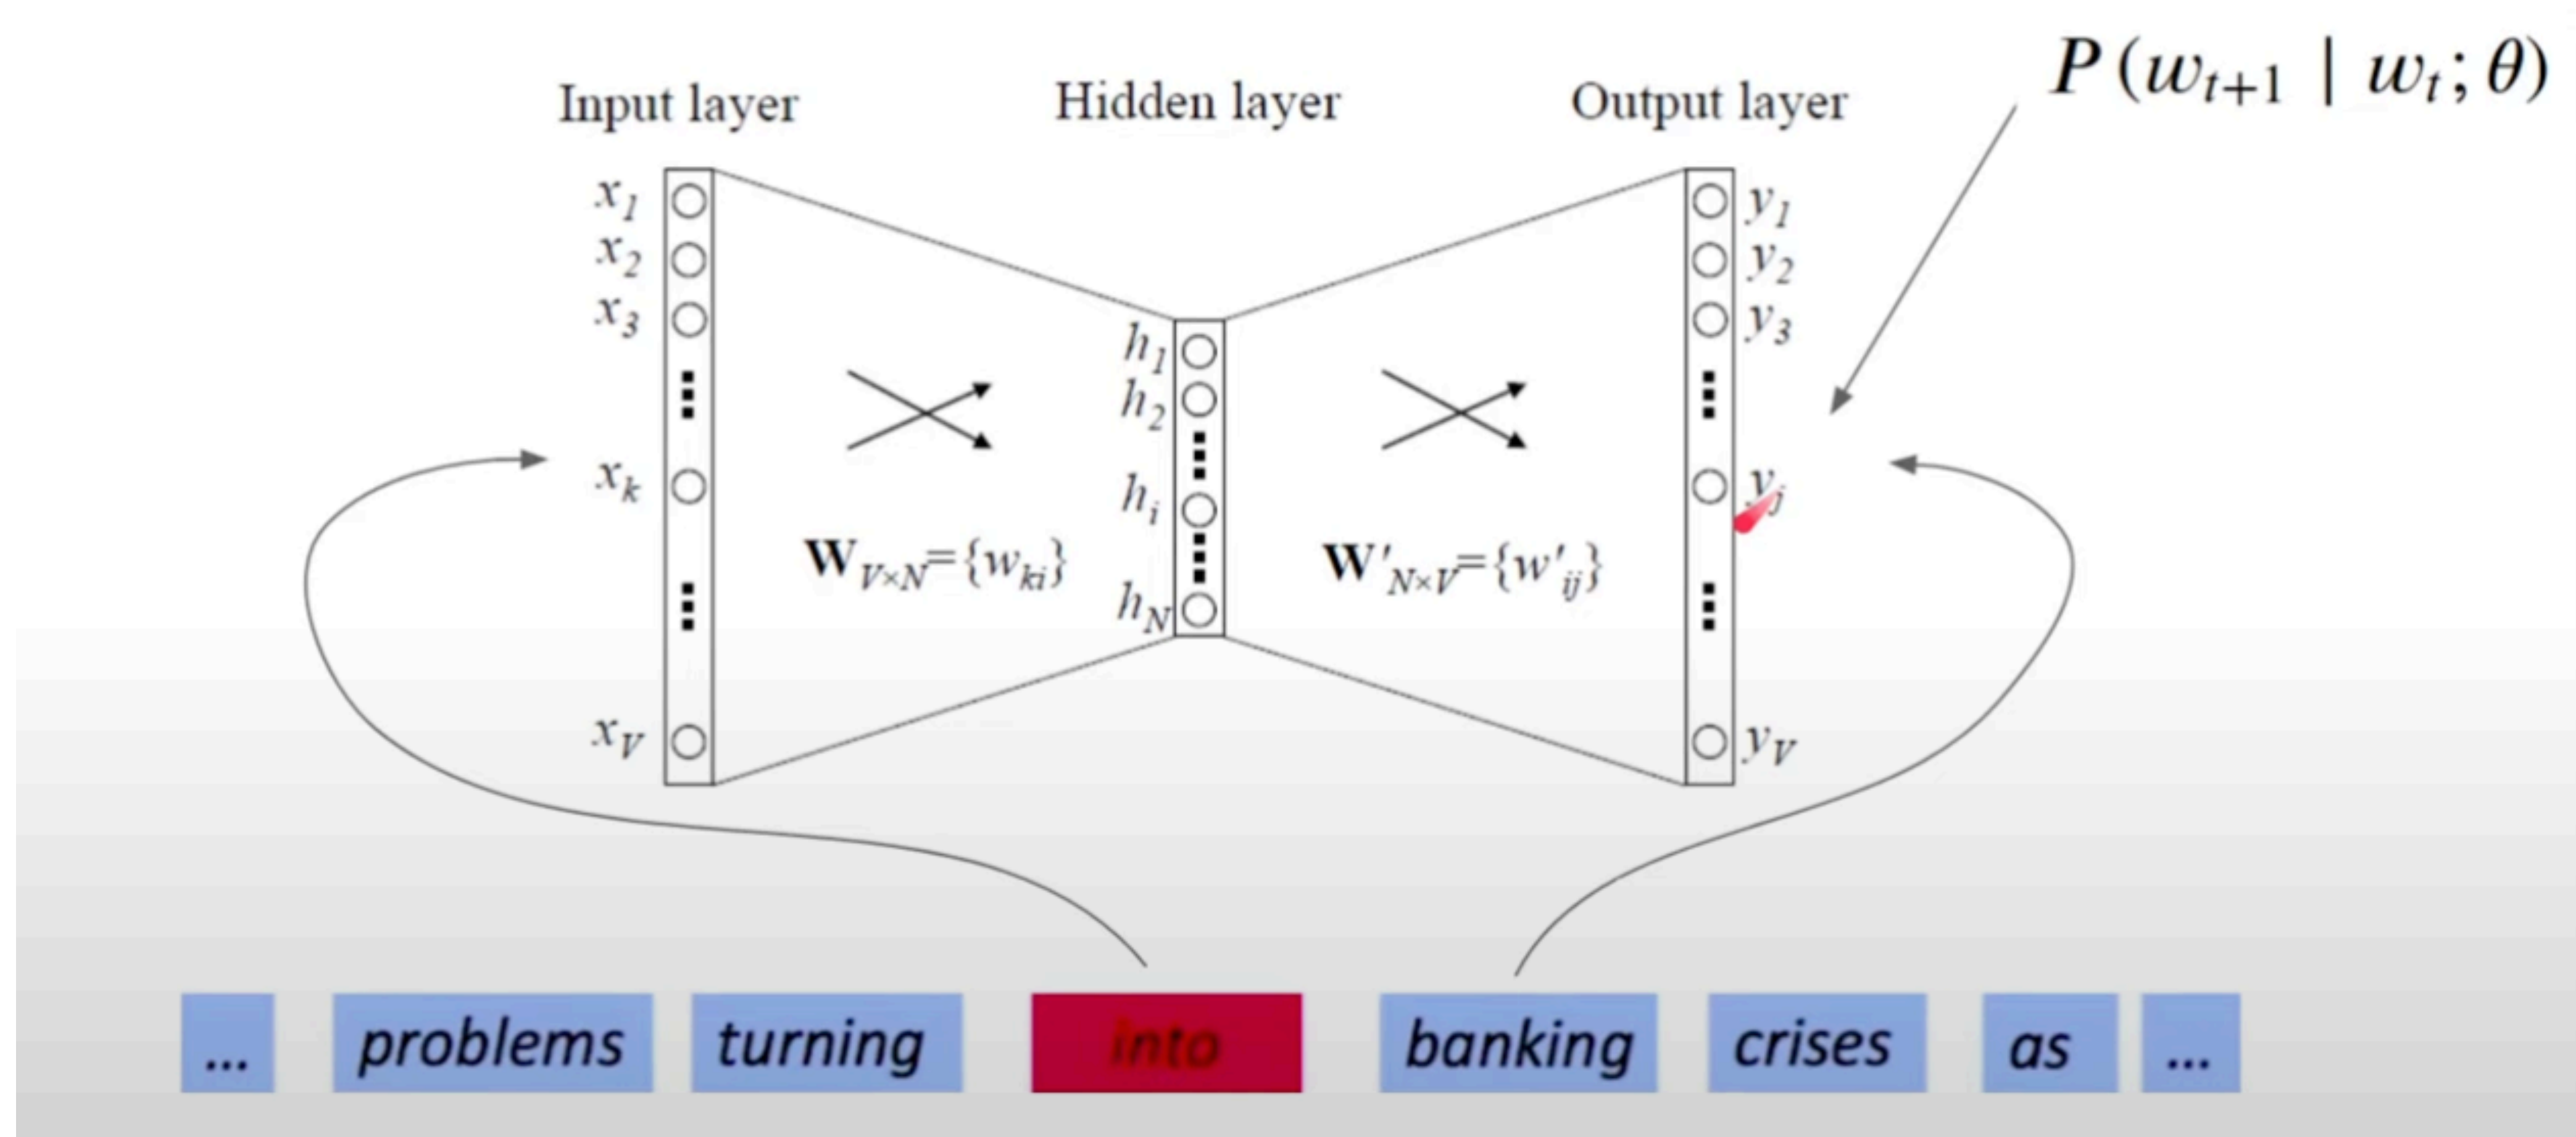
$p(x)$  - распределение, которое мы получили  
 $q(x)$  - распределение истинное




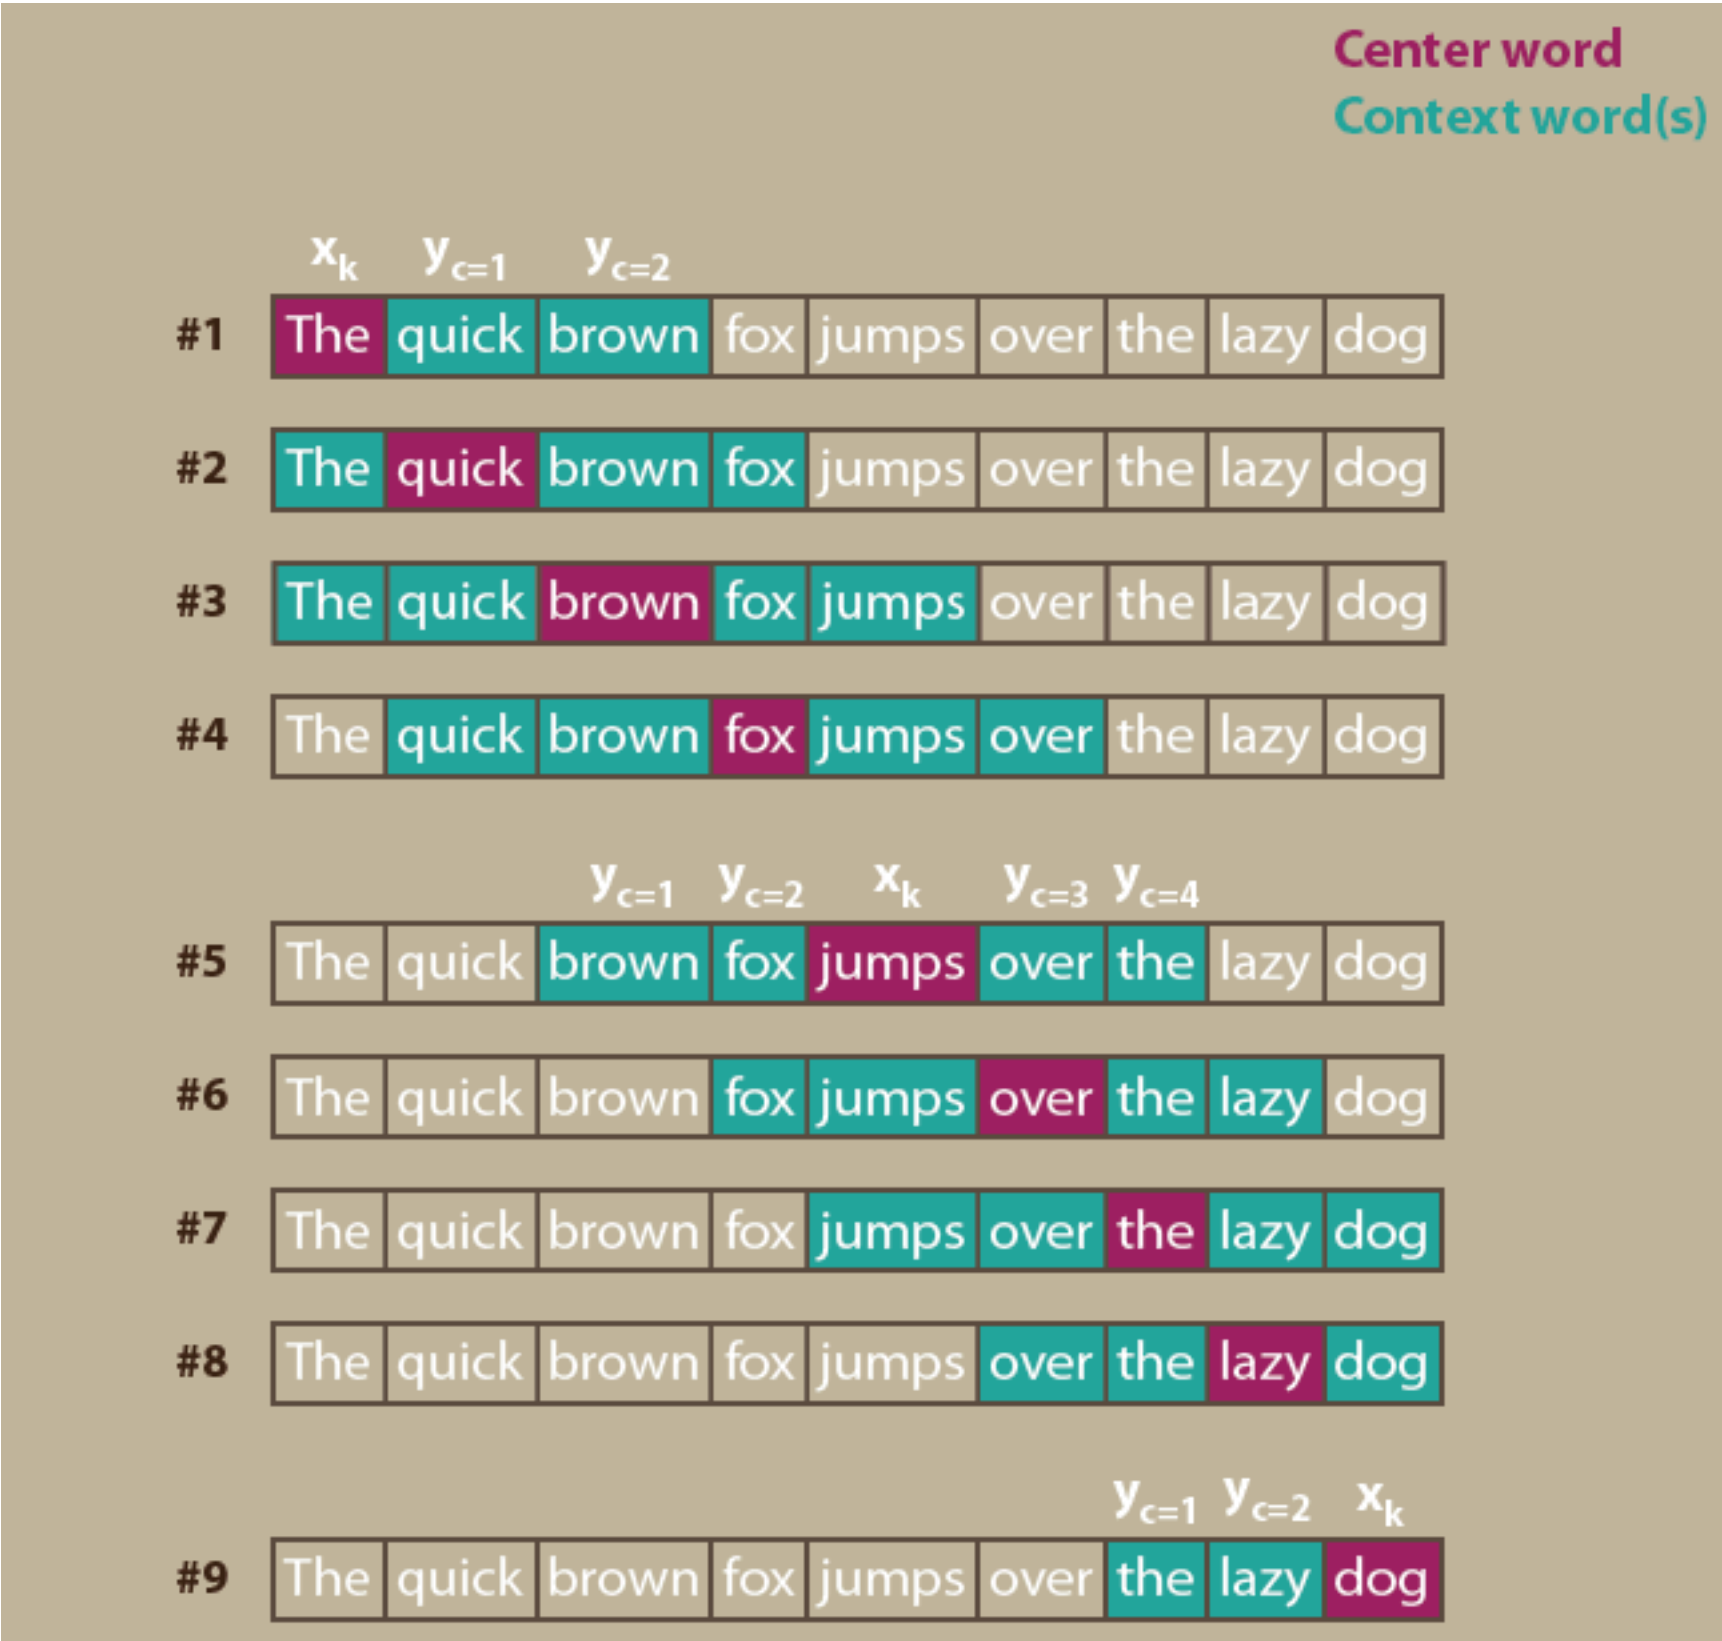
$$H(q, p) = - \sum_x q(x) \log p(x)$$



# Word2vec



# Процесс обучения



# Процесс обучения

$V$  - размер словаря

$N$  - размер эмбединга (200)

$$\mathbf{W} - V \times N$$


$$\mathbf{W}' - N \times V$$

$$\mathbf{W}^T \cdot \mathbf{x} = h \implies (N \times V) \cdot (V \times 1) = N \times 1$$

$$\mathbf{W}^T = [w_1^T w_2^T \dots w_V^T] \Rightarrow \sum_{i=1}^V w_i^T x_i = w_k^T = h$$

$$\mathbf{W}'^T \cdot h = y \implies (V \times N) \cdot (N \times 1) = V \times 1$$


$$\mathbf{W}' = [w'_1 w'_2 \dots w'_V] \Rightarrow y_j = (\mathbf{W}'^T \cdot h)_j = (w'_j)^T w_k^T = \langle w_k, w'_j \rangle$$



Матрица с представлением слова как центрального

Матрица с представлением слова как контекстного

# Процесс обучения



Я люблю пить кофе по утрам одна  $V = 100$

$K = 23$ , позиция слова кофе в словаре,  $j = 45 = \text{пить}$ ,  $[0, 0, 0, 0, 0, 0, \dots, 1, \dots]$

1) вытаскиваем 23 эмбединг, умножаем на матрицу контекстных эмбедингов и получаем распределение на все контекст слова

# Questions

- Какая модель работает быстрее?
- Какую проблему видите?



# Subsampling

**Problem? Some words are not meaningful.**

Each word  $w$  in the training set is discarded with the probability computed by the formula

$$P(w_i) = 1 - \sqrt{\frac{t}{f(w_i)}}$$

t- chosen threshold, words with a frequency greater than t are discarded  
f(w<sub>i</sub>) - frequency of w<sub>i</sub>

*She drinks neither a cup of coffee nor a cup of tea for breakfast.*

*(a cup of + coffee + nor a cup)*

*(drinks neither cup + coffee + cup tea breakfast)*

# Negative sampling

1. Пить и кофе - это близкие слова = контекстные 1
2. Кофе и машина - это разные далекие = не контекстные 0

Problem? Computationally intense to train over the whole vocab.


Dot product of  $v_{cat}$ :

- with  $u_{cute}$  - increase,
- with all other  $u$  - decrease



Dot product of  $v_{cat}$ :

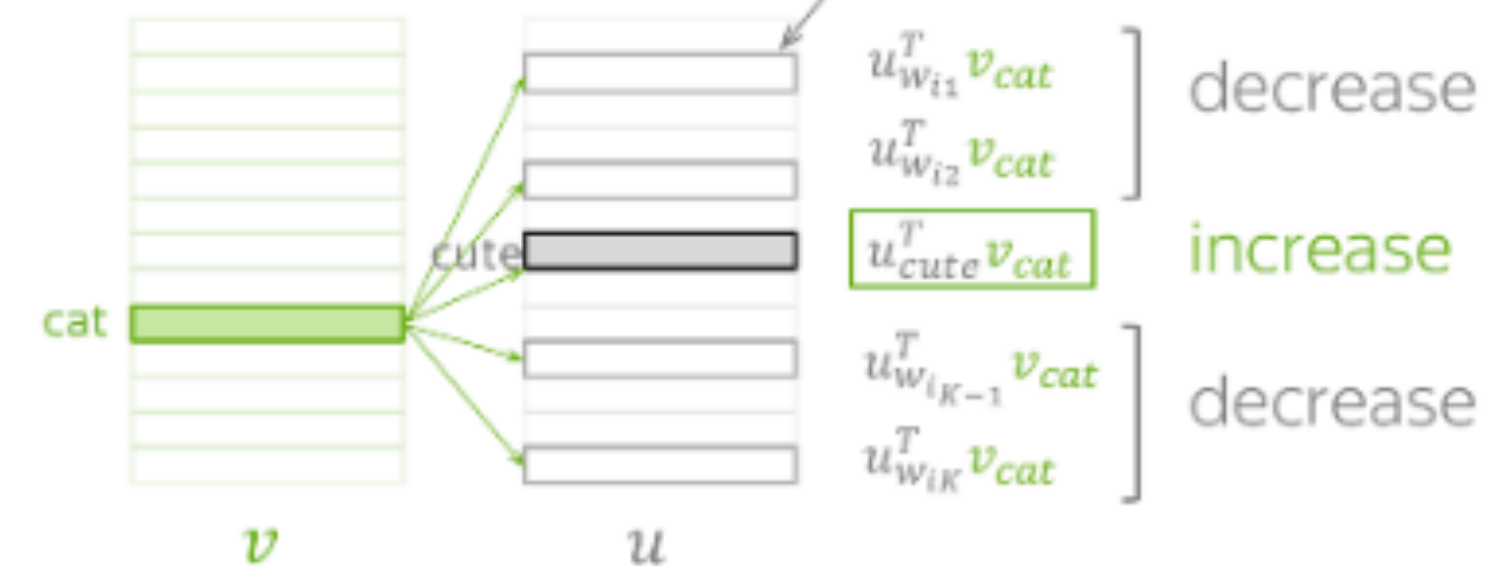
- with  $u_{cute}$  - increase,
- with a subset of other  $u$  - decrease



Parameters to be updated:

- $v_{cat}$
  - $u_w$  for all  $w$  in the vocabulary
- $|V| + 1$  vectors

Negative samples: randomly selected K words



Parameters to be updated:

- $v_{cat}$
  - $u_{cute}$  and  $u_w$  for  $w$  in K negative examples
- $K + 2$  vectors

# Recipe

Somewhat standard setting is:

- **Model:** Skip-Gram with negative sampling;
- **Number of negative examples:** for smaller datasets, 15-20; for huge datasets (which are usually used) it can be 2-5.
- **Embedding dimensionality:** frequently used value is 300, but other variants (e.g., 100 or 50) are also possible.

Я люблю пить **кофе**

**Кофе** - это лучший напиток

Моя мама готовит **кофе** по утрам

Есть **кофе** - есть бодрость

{кофе: [1, 2, 4 23, 424 ],  
пить: [1, 3452423, 4236 ,]}



# Types

# Glove

Problem? Need global info.

Global information from corpus to **learn vectors**

Уменьшаем лосс слов, которые редко встречаются

Before training count occurrences of pairs [word<sub>i</sub> , word<sub>j</sub>] in corpus

Compute probabilities:  $P_{ij} = \frac{Count(v_i, v_j)}{Count(v_i)}, Count(v_i) = \sum_k Count(v_i, v_k)$

Objective function:

$$J(\theta) = \frac{1}{2} \sum_{i,j=1}^W \boxed{f(P_{ij})} (u_i^T v_j - \log P_{ij})^2$$

Discount factor for rare words

Я люблю пить **кофе**

**Кофе** - это лучший напиток

Моя мама готовит **кофе** по утрам

Есть **кофе** - есть бодрость

На самолете кофе никогда не наливают

# FastText

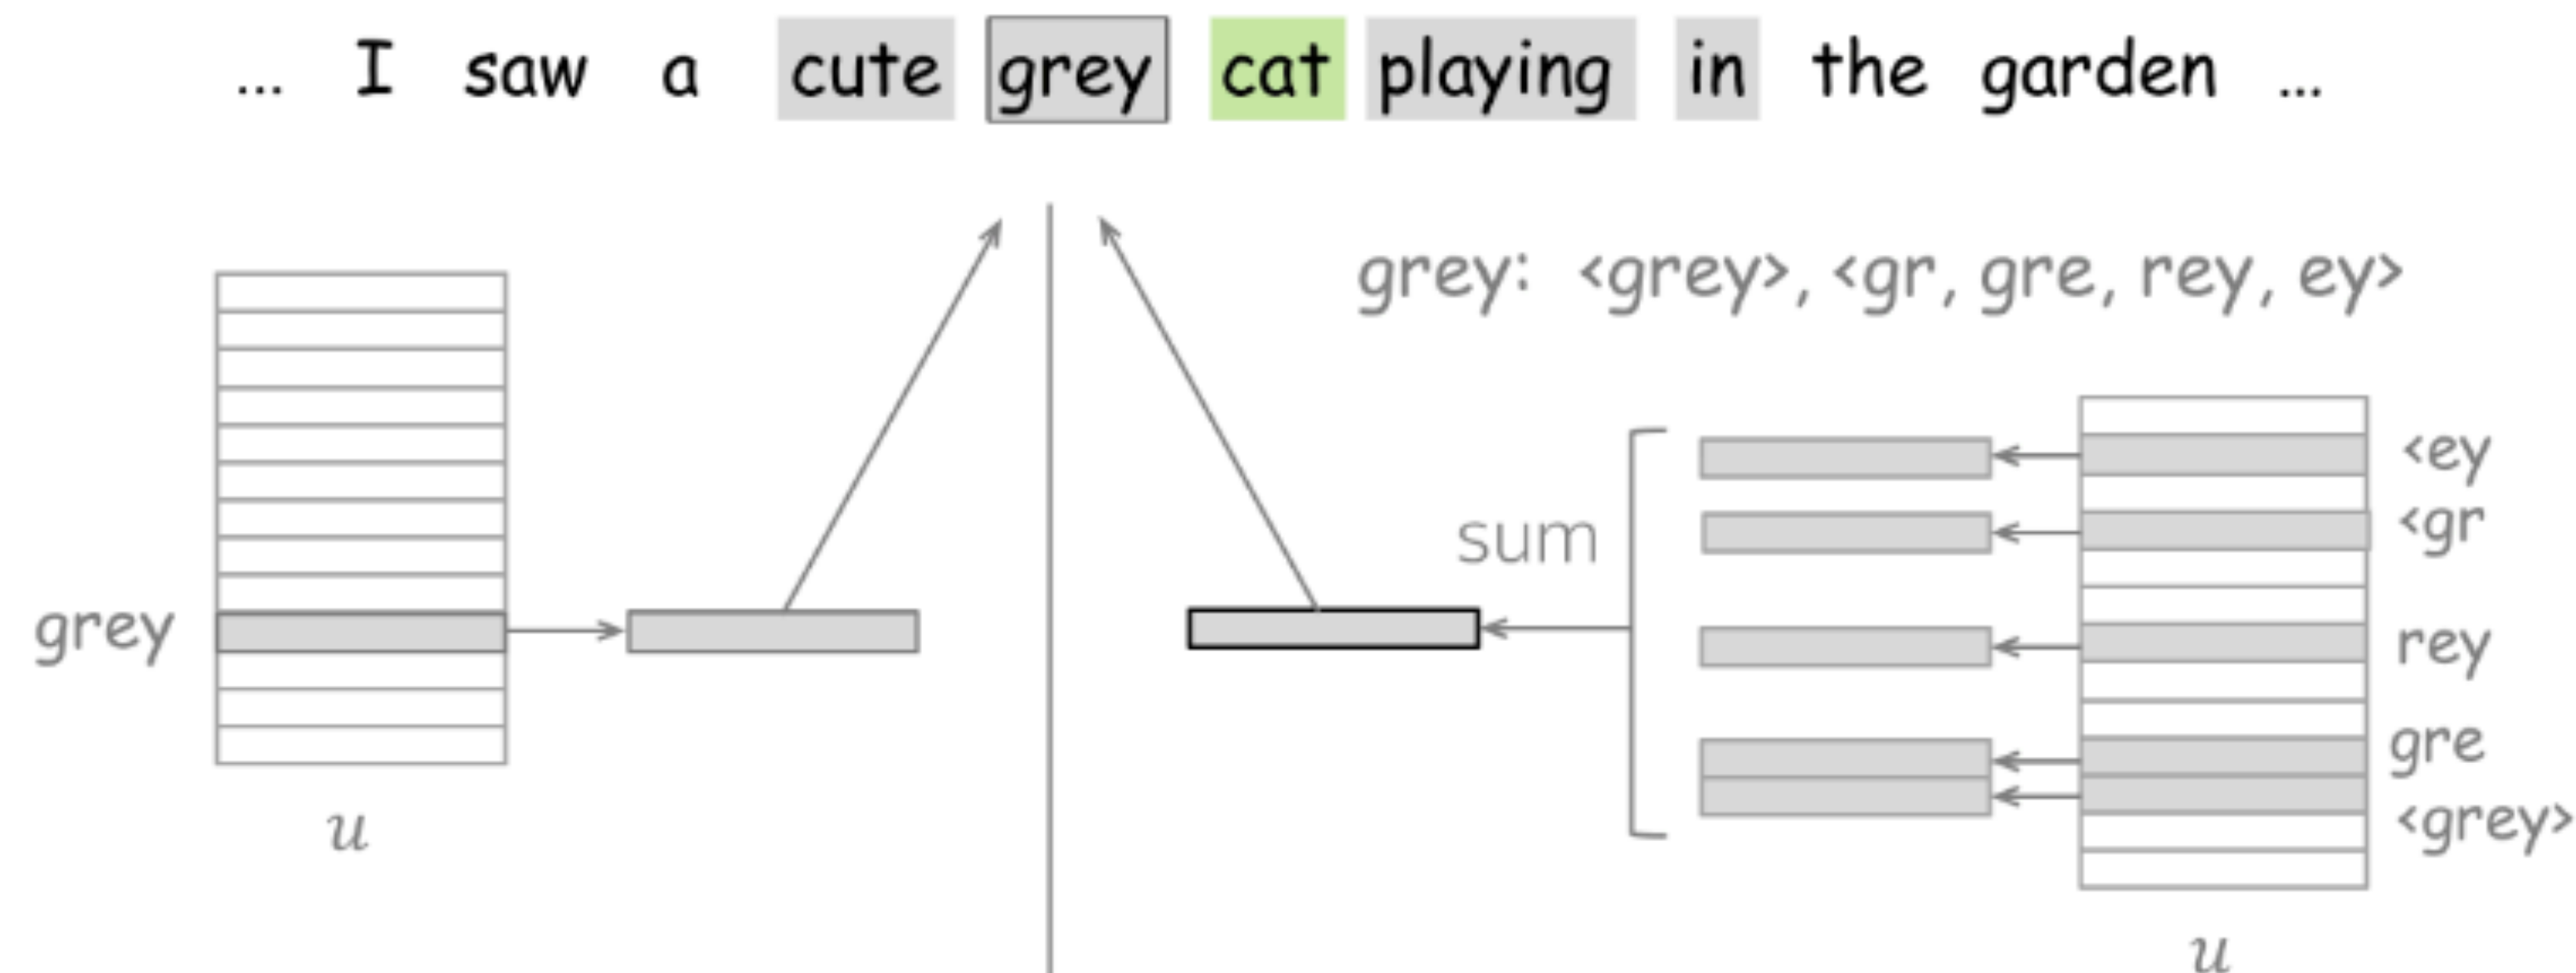
## Problem? Need morph info

Убежал 2  
Прибежал 1  
Сбегал 3  
Бегать 5  
Перебежал 10

Убежал = у + бегать  
Прибежал = При + бегать  
Сбегал = с + бегать

21 обновим эмбединг для бегать

**Out of vocabulary**



### Word2Vec

Vocabulary consists of:

- words

Word vector is:

- one vector from the look-up table

### FastText

Vocabulary consists of:

- words and character n-grams

Word vector is:

- sum of word vector and vectors for its n-grams

# FastText

## Какие есть в этом плюсы?

- better understanding of morphology

By assigning a distinct vector to each word, we ignore morphology.

Giving information about subwords can let the model know that different tokens can be forms of the same word.

- representations for unknown words

Usually, we can represent only those words, which are present in the vocabulary.

Giving information about subwords can help to represent out-of-vocabulary words relying on their spelling.

- handling misspellings

Even if one character in a word is wrong, this is another token, and, therefore, a completely different embedding (or even unknown word). With information about subwords, misspelled word would still be similar to the original one.

# FastText

Какие есть в этом плюсы?



# ELMO

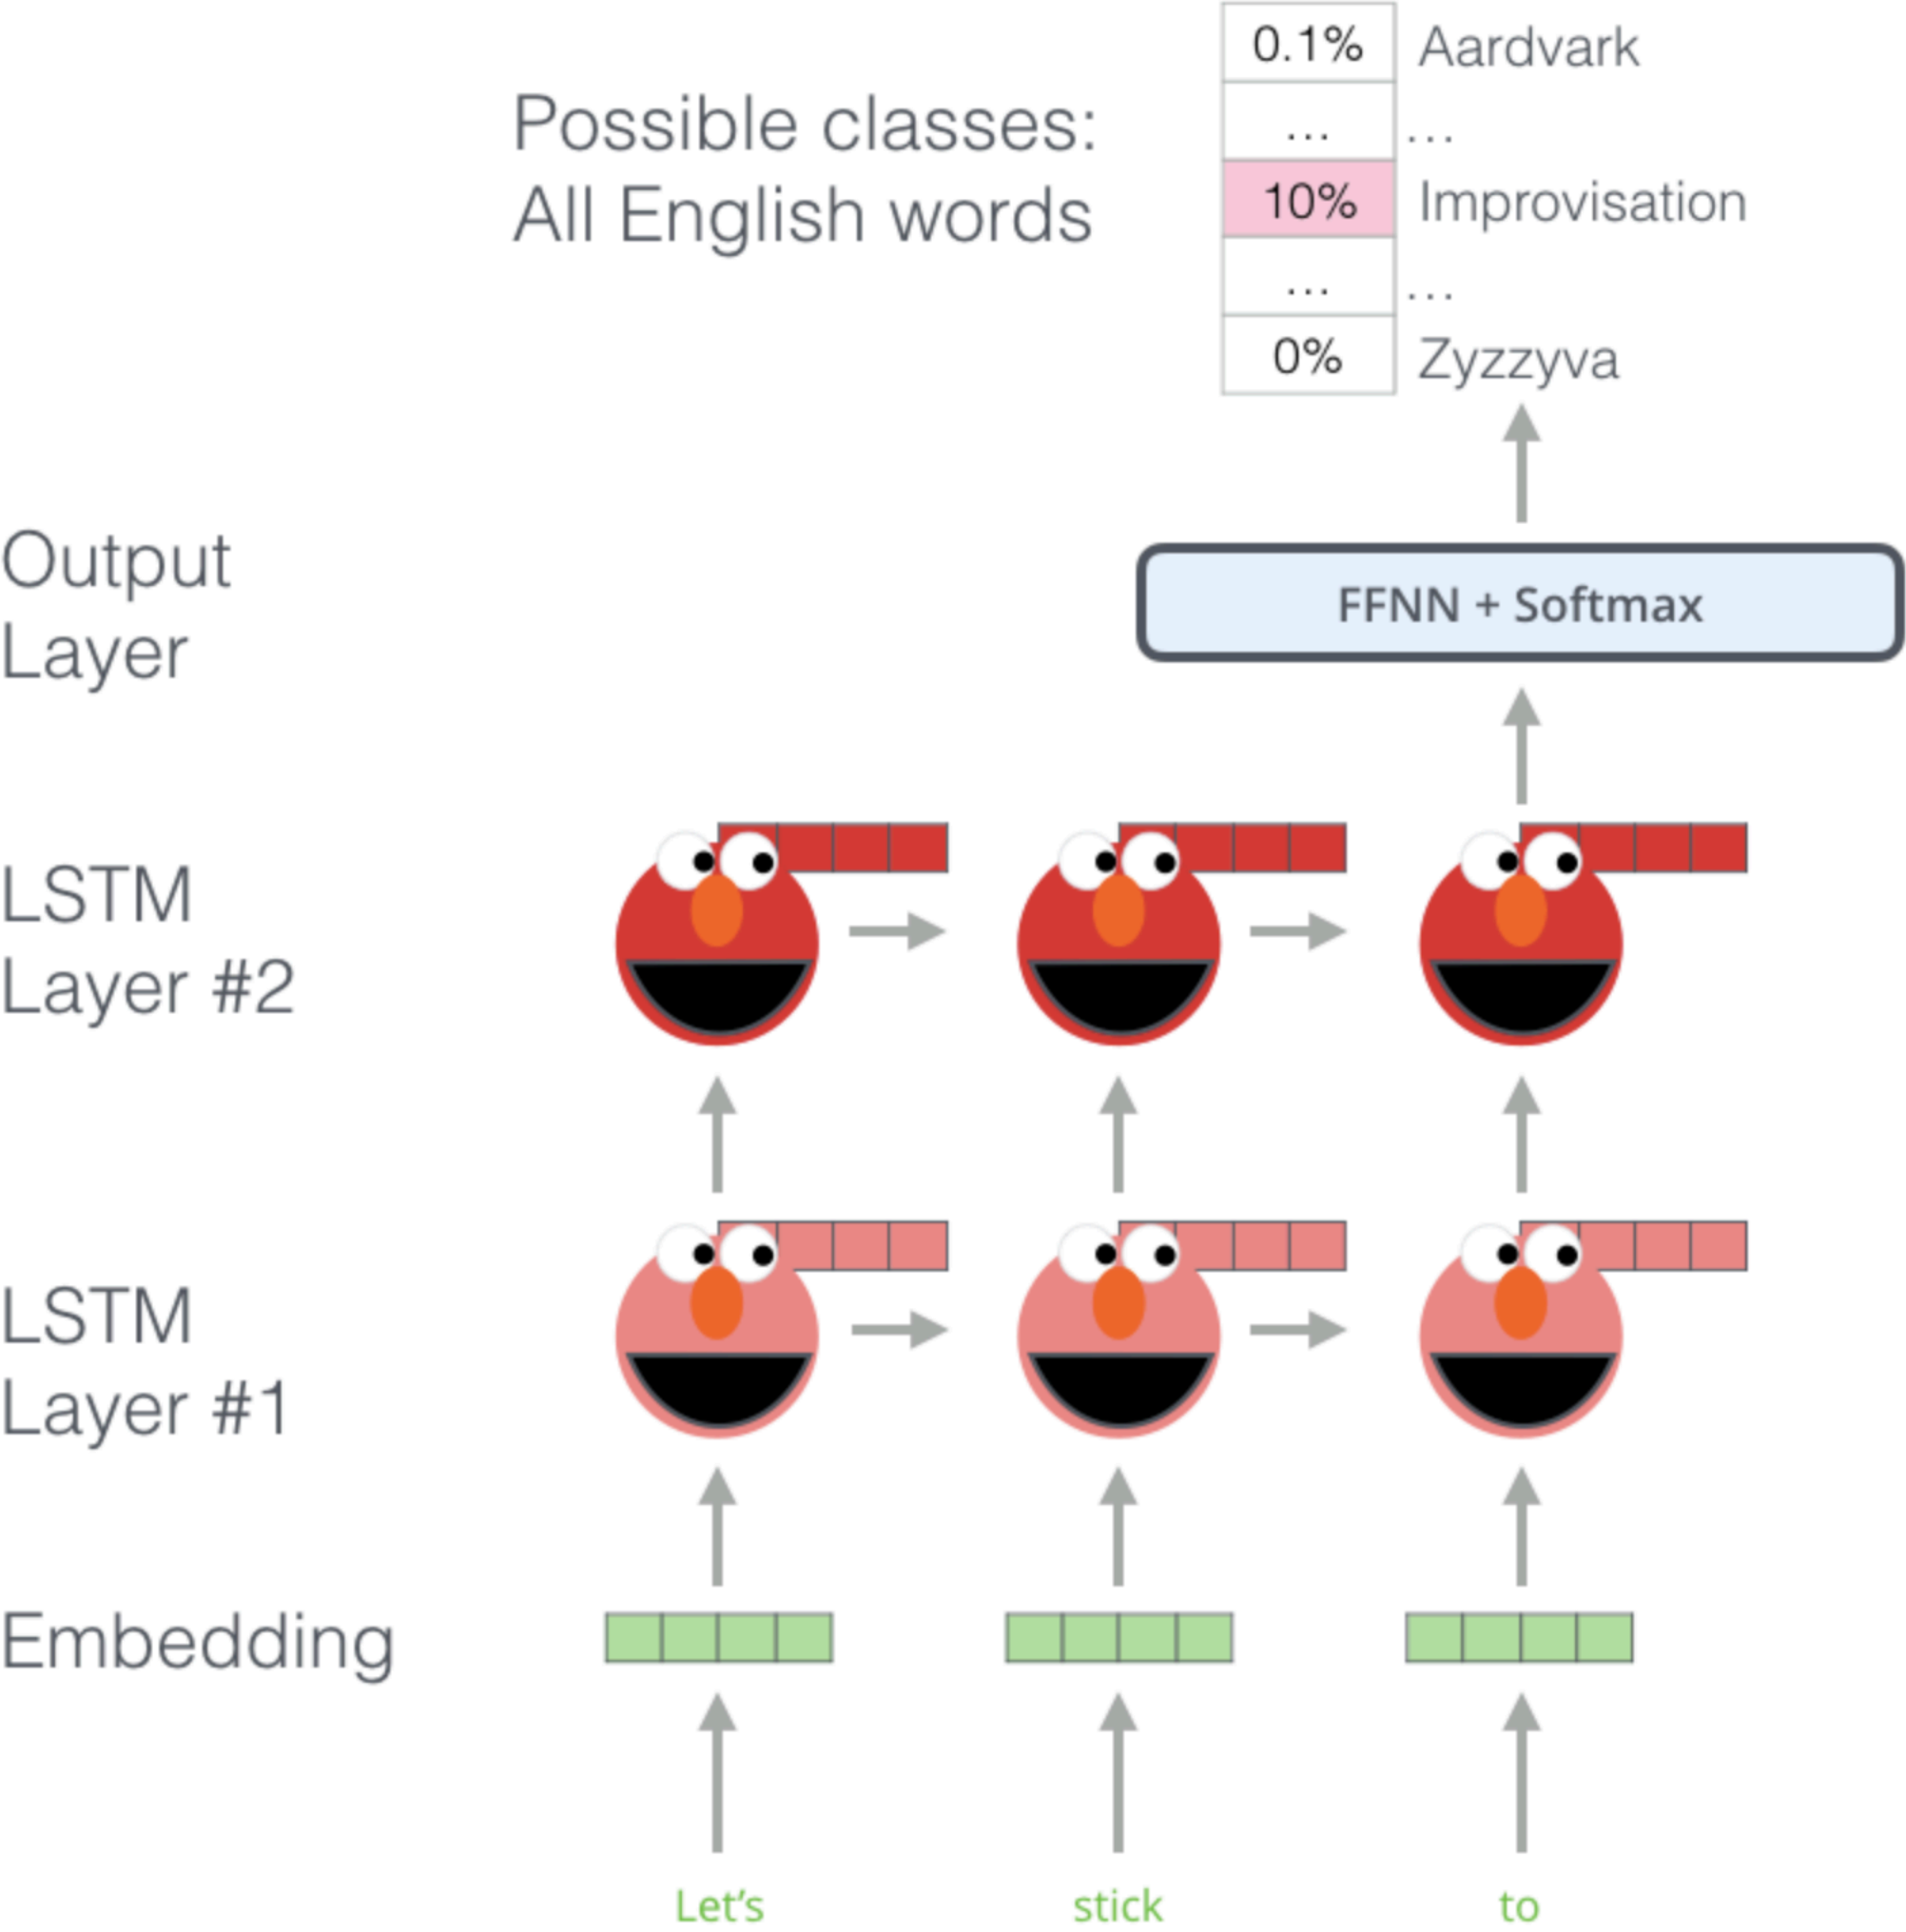
Problem? We need context



# ELMO

## Embeddings from Language Models

Language Modeling task - predict the next word



# ELMO

## Problem? We need context


- \* Language Modeling task - predict the next word
- \* Bidirectional LSTM
- \* Token embeddings

$$ELMo_k^{task} = \gamma_k \cdot (s_0^{task} \cdot x_k + s_1^{task} \cdot h_{1,k} + s_2^{task} \cdot h_{2,k})$$


ELMo comes up with the contextualized embedding through grouping together the hidden states (and initial embedding) in a certain way (concatenation followed by weighted summation).

### Embedding of “stick” in “Let’s stick to” - Step #2

1- Concatenate hidden layers



2- Multiply each vector by a weight based on the task




3- Sum the (now weighted) vectors



ELMo embedding of “stick” for this task in this context

Forward Language Model



Backward Language Model

